

Detect credit card fraud transactions using Logistic Regression

About the dataset

- The dataset contains transactions made by credit cards in September 2013 by European cardholders. This dataset presents transactions that occurred in two days, where we have 492 frauds out of 284,807 transactions. The dataset is highly unbalanced, the positive class (frauds) account for 0.172% of all transactions

```
# importing required libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

import warnings
warnings.filterwarnings('ignore')

# Model selection and training
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import
classification_report, accuracy_score, confusion_matrix
#Threshold
from sklearn.metrics import roc_auc_score
from sklearn.metrics import roc_curve

# importing dataset
df = pd.read_csv('creditcard.csv')

df.shape

(284807, 31)

df.head()
```

	Time	V1	V2	V3	V4	V5	V6
V7 \							
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388
0.239599							
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361
0.078803							
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499
0.791461							
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203
0.237609							

```
4  2.0 -1.158233  0.877737  1.548718  0.403034 -0.407193  0.095921
0.592941
```

```

      V8      V9  ...      V21      V22      V23      V24
V25 \
0  0.098698  0.363787  ... -0.018307  0.277838 -0.110474  0.066928
0.128539
1  0.085102 -0.255425  ... -0.225775 -0.638672  0.101288 -0.339846
0.167170
2  0.247676 -1.514654  ...  0.247998  0.771679  0.909412 -0.689281 -
0.327642
3  0.377436 -1.387024  ... -0.108300  0.005274 -0.190321 -1.175575
0.647376
4 -0.270533  0.817739  ... -0.009431  0.798278 -0.137458  0.141267 -
0.206010
```

```

      V26      V27      V28  Amount  Class
0 -0.189115  0.133558 -0.021053  149.62      0
1  0.125895 -0.008983  0.014724   2.69      0
2 -0.139097 -0.055353 -0.059752  378.66      0
3 -0.221929  0.062723  0.061458  123.50      0
4  0.502292  0.219422  0.215153   69.99      0
```

```
[5 rows x 31 columns]
```

```
df.info() # no need to deal with null values
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Time        284807 non-null  float64
1   V1          284807 non-null  float64
2   V2          284807 non-null  float64
3   V3          284807 non-null  float64
4   V4          284807 non-null  float64
5   V5          284807 non-null  float64
6   V6          284807 non-null  float64
7   V7          284807 non-null  float64
8   V8          284807 non-null  float64
9   V9          284807 non-null  float64
10  V10         284807 non-null  float64
11  V11         284807 non-null  float64
12  V12         284807 non-null  float64
13  V13         284807 non-null  float64
14  V14         284807 non-null  float64
15  V15         284807 non-null  float64
16  V16         284807 non-null  float64
17  V17         284807 non-null  float64
```

```

18 V18      284807 non-null float64
19 V19      284807 non-null float64
20 V20      284807 non-null float64
21 V21      284807 non-null float64
22 V22      284807 non-null float64
23 V23      284807 non-null float64
24 V24      284807 non-null float64
25 V25      284807 non-null float64
26 V26      284807 non-null float64
27 V27      284807 non-null float64
28 V28      284807 non-null float64
29 Amount   284807 non-null float64
30 Class    284807 non-null int64

```

dtypes: float64(30), int64(1)

memory usage: 67.4 MB

df.describe()

	Time	V1	V2	V3
V4 \				
count	284807.000000	2.848070e+05	2.848070e+05	2.848070e+05
mean	94813.859575	1.168375e-15	3.416908e-16	-1.379537e-15
std	47488.145955	1.958696e+00	1.651309e+00	1.516255e+00
min	0.000000	-5.640751e+01	-7.271573e+01	-4.832559e+01
25%	54201.500000	-9.203734e-01	-5.985499e-01	-8.903648e-01
50%	84692.000000	1.810880e-02	6.548556e-02	1.798463e-01
75%	139320.500000	1.315642e+00	8.037239e-01	1.027196e+00
max	172792.000000	2.454930e+00	2.205773e+01	9.382558e+00

	V5	V6	V7	V8
V9 \				
count	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05
mean	9.604066e-16	1.487313e-15	-5.556467e-16	1.213481e-16
std	1.380247e+00	1.332271e+00	1.237094e+00	1.194353e+00
min	-1.137433e+02	-2.616051e+01	-4.355724e+01	-7.321672e+01
25%	-6.915971e-01	-7.682956e-01	-5.540759e-01	-2.086297e-01
50%	-5.433583e-02	-2.741871e-01	4.010308e-02	2.235804e-02

```

5.142873e-02
75%      6.119264e-01   3.985649e-01   5.704361e-01   3.273459e-01
5.971390e-01
max       3.480167e+01   7.330163e+01   1.205895e+02   2.000721e+01
1.559499e+01

count      ...      V21      V22      V23      V24  \
mean      ...      1.654067e-16 -3.568593e-16  2.578648e-16  4.473266e-15
std       ...      7.345240e-01  7.257016e-01  6.244603e-01  6.056471e-01
min       ...      -3.483038e+01 -1.093314e+01 -4.480774e+01 -2.836627e+00
25%       ...      -2.283949e-01 -5.423504e-01 -1.618463e-01 -3.545861e-01
50%       ...      -2.945017e-02  6.781943e-03 -1.119293e-02  4.097606e-02
75%       ...      1.863772e-01  5.285536e-01  1.476421e-01  4.395266e-01
max       ...      2.720284e+01  1.050309e+01  2.252841e+01  4.584549e+00

Amount  \
count  2.848070e+05  2.848070e+05  2.848070e+05  2.848070e+05
284807.000000
mean   5.340915e-16  1.683437e-15 -3.660091e-16 -1.227390e-16
88.349619
std    5.212781e-01  4.822270e-01  4.036325e-01  3.300833e-01
250.120109
min    -1.029540e+01 -2.604551e+00 -2.256568e+01 -1.543008e+01
0.000000
25%    -3.171451e-01 -3.269839e-01 -7.083953e-02 -5.295979e-02
5.600000
50%     1.659350e-02 -5.213911e-02  1.342146e-03  1.124383e-02
22.000000
75%     3.507156e-01  2.409522e-01  9.104512e-02  7.827995e-02
77.165000
max     7.519589e+00  3.517346e+00  3.161220e+01  3.384781e+01
25691.160000

count      Class
mean      0.001727
std       0.041527
min       0.000000
25%       0.000000
50%       0.000000
75%       0.000000
max       1.000000

[8 rows x 31 columns]

# Checking for duplicate records
df.duplicated().sum()

```

1081

```
# we found 1081 duplicate records and now we are dropping those
df.drop_duplicates(inplace=True)
print(f'New dataset without duplicates contains {df.shape[0]} rows and
{df.shape[1]} columns')
```

New dataset without duplicates contains 283726 rows and 31 columns


```
# defining colours
colors = sns.color_palette("Set2")
colors
```

```
[ (0.4, 0.7607843137254902, 0.6470588235294118),
  (0.9882352941176471, 0.5529411764705883, 0.3843137254901961),
  (0.5529411764705883, 0.6274509803921569, 0.796078431372549),
  (0.9058823529411765, 0.5411764705882353, 0.7647058823529411),
  (0.6509803921568628, 0.8470588235294118, 0.32941176470588235),
  (1.0, 0.8509803921568627, 0.1843137254901961),
  (0.8980392156862745, 0.7686274509803922, 0.5803921568627451),
  (0.7019607843137254, 0.7019607843137254, 0.7019607843137254)]
```

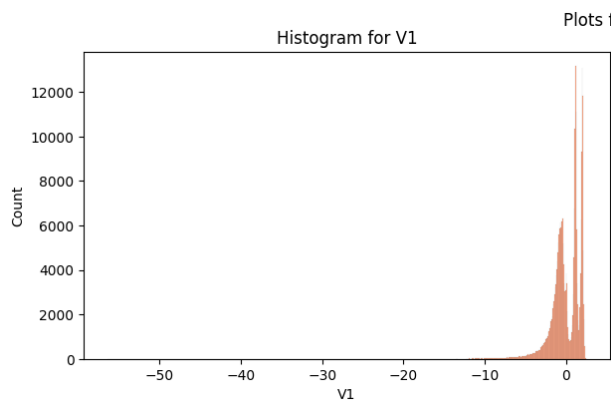
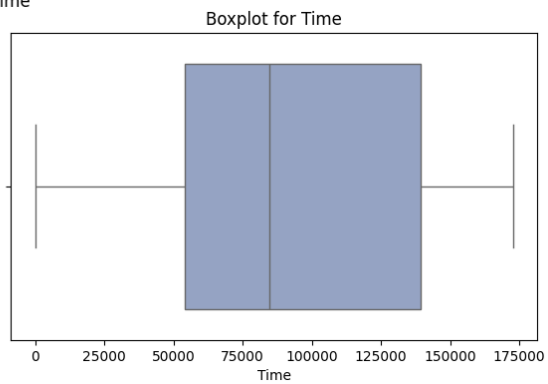
```
# Here all columns are numerical columns so we can directly plot
countplot,boxplot,distplot etc
# Defining a function for representing numerical columns
def numerical_plots(column_name):
    fig,axes = plt.subplots(1,2,figsize = (15,4))
    fig.suptitle(f'Plots for column {column_name}')
    # Histogram
    sns.histplot(data=df,x=column_name,ax=axes[0],color=colors[1])
    axes[0].set_title(f'Histogram for {column_name}')
    # Boxplot
    sns.boxplot(data=df,x=column_name,ax=axes[1],color=colors[2])
    axes[1].set_title(f'Boxplot for {column_name}')

    plt.show()

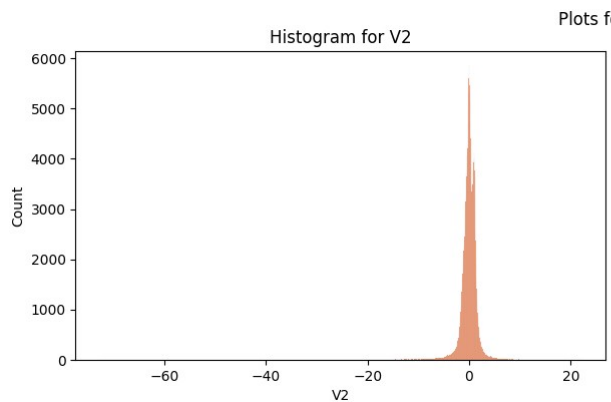

# Plotting histogram and boxplot for all numerical columns
for i in df.columns:
    numerical_plots(i)
```



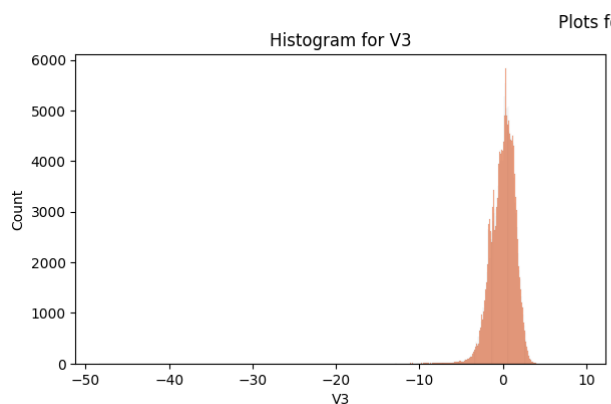

Plots for column Time




Plots for column V1




Plots for column V2





Plots for column V3


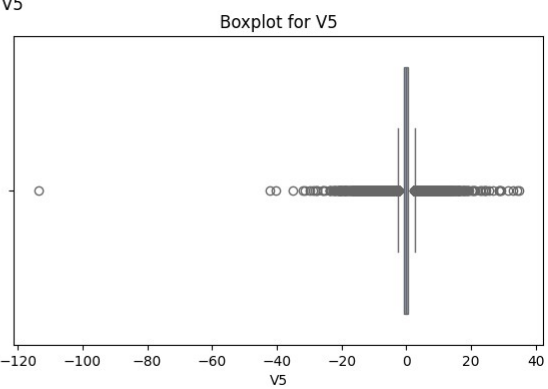




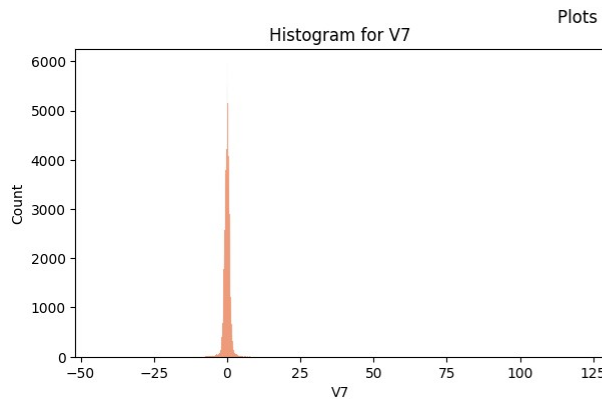

Plots for column V4



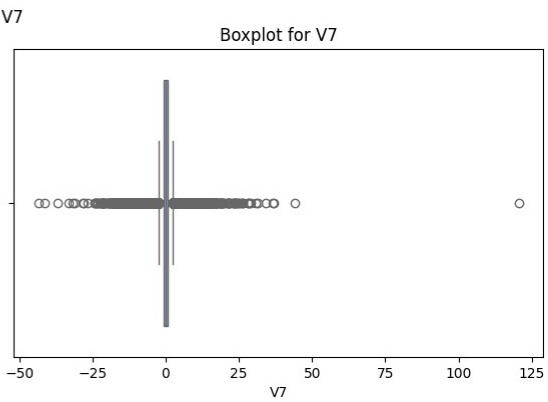
Plots for column V5




Plots for column V6





Plots for column V7

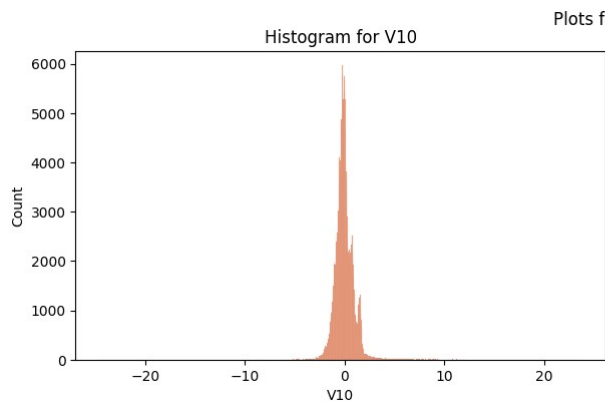






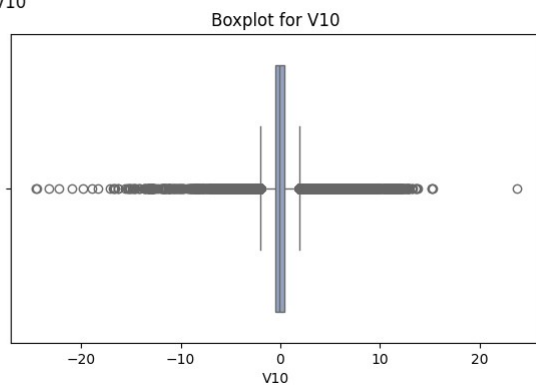
Plots for column V8



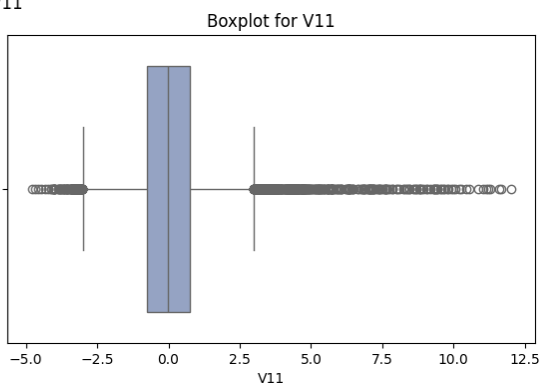
Plots for column V9




Plots for column V10


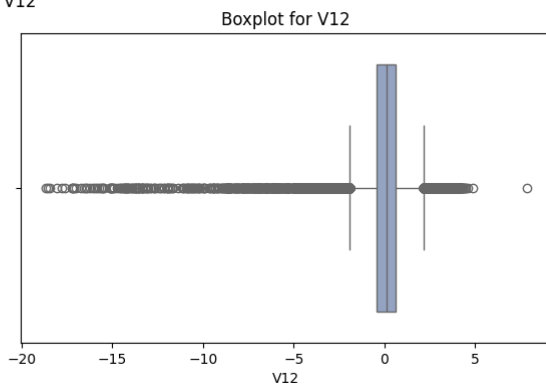


Plots for column V11


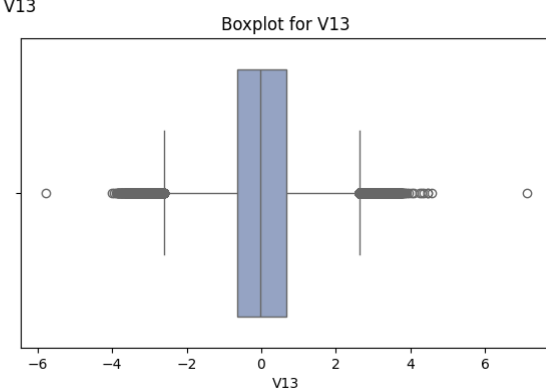






Plots for column V12




Plots for column V13




Plots for column V14





Plots for column V15









Plots for column V16




Plots for column V17





Plots for column V18





Plots for column V19





Plots for column V20





Plots for column V21




Plots for column V22





Plots for column V23









Plots for column V24




Plots for column V25

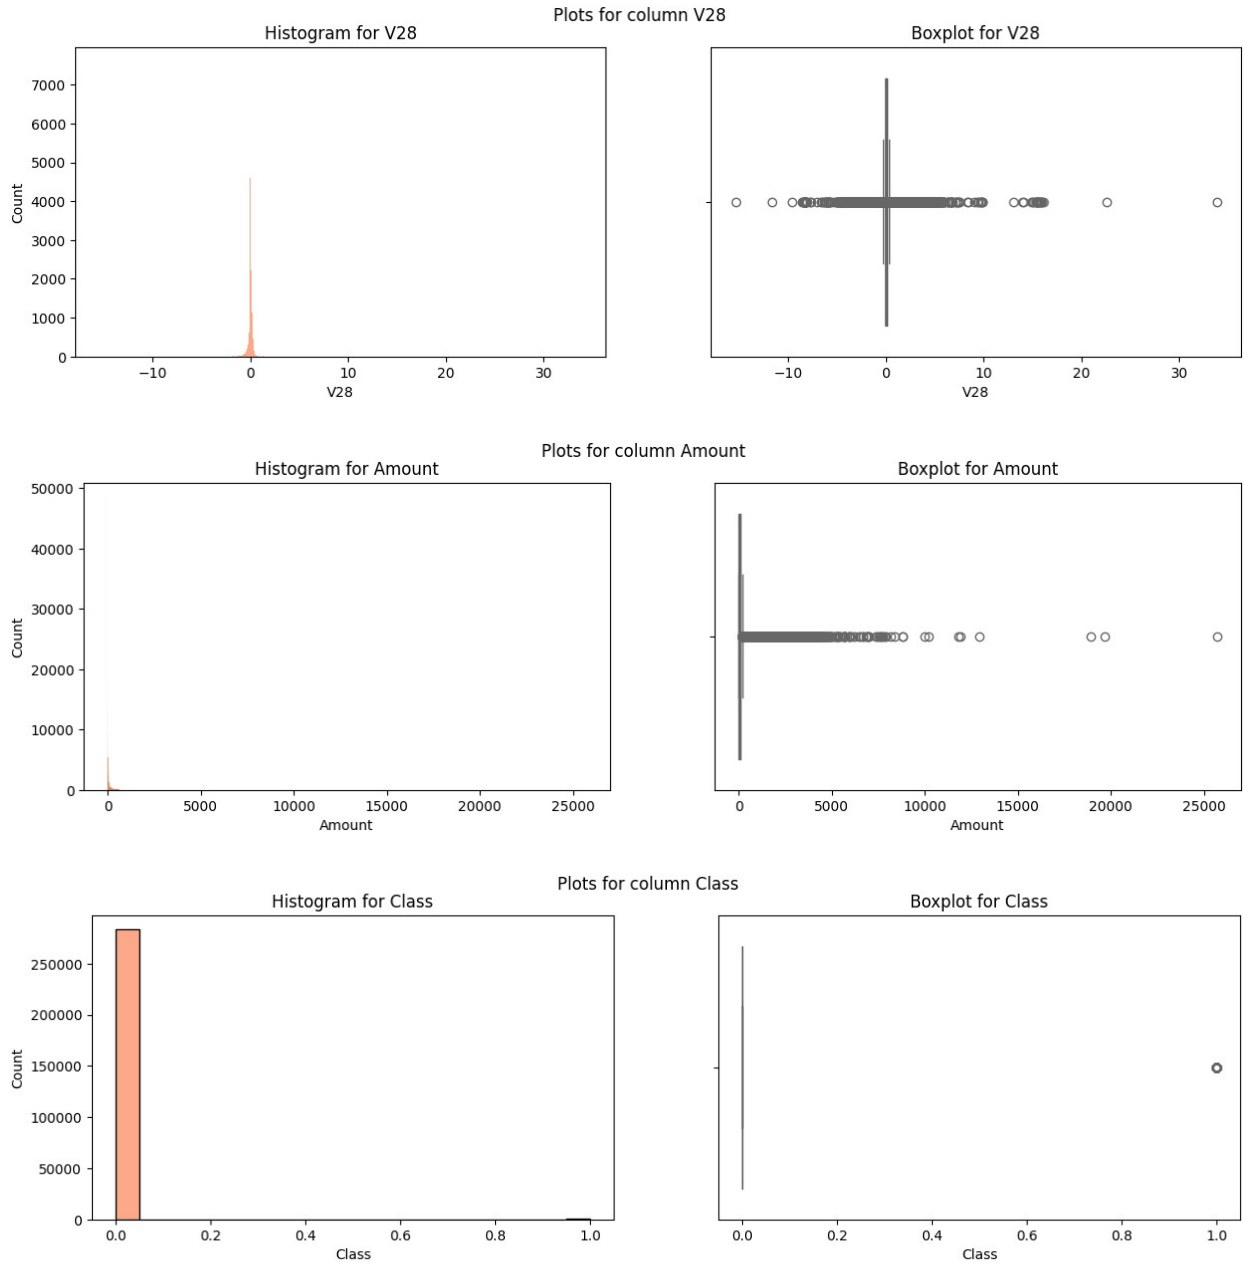


Plots for column V26




Plots for column V27





```
# Checking the correlation between all numerical columns
correlation_matrix = df.corr()
plt.figure(figsize=(14, 12))
sns.heatmap(correlation_matrix, cmap='coolwarm', annot=True,
fmt=".2f", linewidths=0.5)
plt.show()
```



```
# We are scaling data using standard scalar
# splitting data
x_dummy = df.drop(columns=['Class'],axis=1)
Y = df['Class']

#Standard scalar
sc = StandardScaler()
X = sc.fit_transform(x_dummy)

#Splitting data into train and test
x_train,x_test,y_train,y_test =
train_test_split(X,Y,test_size=0.2,random_state=0)
print(f'x_train : {x_train.shape},y_train : {y_train.shape},x_test :
{x_test.shape},y_test : {y_test.shape}')
```

```

x_train : (226980, 30),y_train : (226980,),x_test : (56746, 30),y_test
: (56746,)

# Defining a function for training model and calculating
classification report
def find_logistic_regression(x_train,y_train,x_test):
    # calling logistic regression model and fitting data
    lgr = LogisticRegression()
    lgr.fit(x_train,y_train)
    # prediction for train data
    y_pred_train = lgr.predict(x_train)
    # prediction for test data
    y_pred_test = lgr.predict(x_test)
    # classification report for train model
    y_train_cl_report =
classification_report(y_train,y_pred_train,target_names = ['Non
Fraud','Fraud'])
    # classification report for test model
    y_test_cl_report =
classification_report(y_test,y_pred_test,target_names = ['Non
Fraud','Fraud'])
    # printing the classification report and returning y_test_pred and
logistic regression model
    print('_' * 100)
    print('CLASSIFICATION REPORT FOR TRAIN MODEL')
    print('_' * 100)
    print(y_train_cl_report)
    print('_' * 100)
    print('CLASSIFICATION REPORT FOR TEST MODEL')
    print('_' * 100)
    print(y_test_cl_report)
    return y_pred_test,lgr

# calling function for given train and test data
y_pred,lgr = find_logistic_regression(x_train,y_train,x_test)

```

CLASSIFICATION REPORT FOR TRAIN MODEL

	precision	recall	f1-score	support
Non Fraud	1.00	1.00	1.00	226594
Fraud	0.89	0.63	0.74	386
accuracy			1.00	226980
macro avg	0.94	0.81	0.87	226980
weighted avg	1.00	1.00	1.00	226980

CLASSIFICATION REPORT FOR TEST MODEL


	precision	recall	f1-score	support
Non Fraud	1.00	1.00	1.00	56659
Fraud	0.83	0.55	0.66	87
accuracy			1.00	56746
macro avg	0.91	0.78	0.83	56746
weighted avg	1.00	1.00	1.00	56746

Defining a custom function for confusion matrix

```
def plot_confusion_matrix(y_test,y_pred):  
    cnfn_matrx = confusion_matrix(y_test,y_pred)  
    labels = ['Non Fraud','Fraud']  
  
    sns.heatmap(cnfn_matrx,annot=True,fmt='d',xticklabels=labels,yticklabels=labels)  
    plt.xlabel('Actual')  
    plt.ylabel('Predicted')  
    plt.show()
```

plotting confusion matrix for above logistic regression model

```
plot_confusion_matrix(y_test,y_pred)
```

```
df['Class'].value_counts()
Class
0    283253
1      473
Name: count, dtype: int64

# Different df for both groups
df1 = df[df['Class'] == 0].sample(n = 473, random_state = 42)
df2 = df[df['Class'] == 1].sample(n = 473, random_state = 42)
print(f'Non fraud df with shape {df1.shape} and Fraud df with shape {df2.shape}')
```

Non fraud df with shape (473, 31) and Fraud df with shape (473, 31)

```
#combining them
df_final = pd.concat([df1, df2], ignore_index=True)
df_final['Class'].value_counts()
Class
0      473
1      473
Name: count, dtype: int64
```

```

# splitting for x and y
X_balanced_dummy = df_final.drop(columns=['Class'],axis=1)
Y_balanced = df_final['Class']

# Scaling independent variables using Standard Scalar
sc_bal = StandardScaler()
X_balanced = sc_bal.fit_transform(X_balanced_dummy)

# Splitting for train and test
x_train_bal,x_test_bal,y_train_bal,y_test_bal =
train_test_split(X_balanced,Y_balanced,test_size=0.2,random_state=0)
print(f'x_train : {x_train_bal.shape},y_train :
{y_train_bal.shape},x_test : {x_test_bal.shape},y_test :
{y_test_bal.shape}')

x_train : (756, 30),y_train : (756,),x_test : (190, 30),y_test :
(190,)

# training model and printing classification report
lgr = LogisticRegression()
lgr.fit(x_train_bal,y_train_bal)
# prediction for train data
y_pred_train_bal = lgr.predict(x_train_bal)
# prediction for test data
y_pred_test_bal = lgr.predict(x_test_bal)
# classification report for train model
y_train_cl_report =
classification_report(y_train_bal,y_pred_train_bal)
# classification report for test model
y_test_cl_report = classification_report(y_test_bal,y_pred_test_bal)
# printing the classification report and returning y_test_pred and
logistic regression model
print('_' * 100)
print('CLASSIFICATION REPORT FOR TRAIN MODEL')
print('_' * 100)
print(y_train_cl_report)
print('_' * 100)
print('CLASSIFICATION REPORT FOR TEST MODEL')
print('_' * 100)
print(y_test_cl_report)

```

CLASSIFICATION REPORT FOR TRAIN MODEL

	precision	recall	f1-score	support
0	0.92	0.98	0.95	378
1	0.98	0.92	0.95	378


accuracy			0.95	756
macro avg	0.95	0.95	0.95	756
weighted avg	0.95	0.95	0.95	756

CLASSIFICATION REPORT FOR TEST MODEL

	precision	recall	f1-score	support
0	0.93	0.97	0.95	95
1	0.97	0.93	0.95	95

accuracy			0.95	190
macro avg	0.95	0.95	0.95	190
weighted avg	0.95	0.95	0.95	190

```
# plotting confusion matrix
plot_confusion_matrix(y_test_bal,y_pred_test_bal)
```



Thank you